

Machine learning for vision and multimedia

(01URPOV)

Preliminaries – Vectorization in Python and Numpy
Lia Morra

2024 – 2025

Why vectorization



1 (car) or 0 (non car)?

“vectorization”

$$X = \begin{bmatrix} | & | & | & | & | \\ \mathbf{x}^{(1)} & \mathbf{x}^{(2)} & \dots & \mathbf{x}^{(m)} & \\ | & | & | & | & | \end{bmatrix}$$

N, or n_x features

m training examples

$$Y = \begin{bmatrix} y^{(1)} & y^{(2)} & \dots & y^{(m)} \end{bmatrix}$$

m

Why vectorization

- Machine/deep learning libraries operate on **tensors**, a generalization of arrays and matrices to N-dimension
- *“In the general case, an array of numbers arranged on a regular grid with a variable number of axes is known as a tensor” (I. Goodfellow, Deep learning, p.33)*
- **Vectorization** refers to a style of programming in which traditional **for** loops and **if** statements are replaced by vector operations
- Vectorization supports the efficient **Single Instruction Multiple Data (SIMD)** computational model (GPU)

Vectorization: NumPy

- NumPy is the core library for scientific computing in Python
- It provides a high-performance multidimensional array object, and functions/instructions for efficient manipulation
- Many ML libraries (SciPy, scikit-learn) are built on numpy arrays, which are also useful for loading and preprocessing data
- Tensorflow/Pytorch have their own Tensor implementation but many concepts are similar (e.g., broadcasting)

NumPy vs. Matlab


- NumPy is quite similar, in many ways, to Matlab
- Like Matlab, it is heavily optimized for vectorized operations
 - ♦ It is considerable easier to switch from Matlab to Python/numpy than from C/C++/Java or other procedural languages
 - ♦ There are however many differences in notation
- A quick reference of the main differences is available here:
 - ♦ <https://numpy.org/doc/stable/user/numpy-for-matlab-users.html>

NumPy arrays

- An array is a grid of values, all of the same type, indexed by nonnegative integers
- The number of dimensions is the **rank** of the array
- The **shape** of an array is a tuple of integers giving the size of the array along each dimension

NumPy arrays (II)


1D array



axis 0 →

shape: (4,)

2D array




axis 0 ↓

axis 1 →

shape: (2, 3)

3D array



shape: (4, 3, 2)

Creating arrays

- Arrays can be created using a variety of functions:
 - ♦ each initializes the values differently (zeros, ones, eyes, random)
- All take as input the shape
 - ♦ (1, 2, 2) creates a tensor of rank 3 with dimensions equal to 1, 2 and 2 along axis 0, 1 and 2, respectively
 - ♦ (2, 2) creates a tensor of rank 2 with dimensions equal to 2 and 2 along axis 0 and 1, respectively
 - ♦ (2, 256, 256, 3) creates a tensor of rank 4 with dimensions equal to 2, 256, 256 and 3, respectively, or... a batch of 2 images, of size 256 by 256, with 3 channels (RGB)

Indexing and slicing

- **Integer indexing:** tensor values can be accessed by their position (from 0 to $n - 1$)
 - ♦ $X[2]$: retrieves the 3rd value
 - ♦ $X[-2]$: negative indices count from the end of the array
 - ♦ $X[i,j]$: retrieves the value at position (i,j) from a 2-dimensional array
- **Slicing** provides a concise syntax to access subarrays
 - ♦ $X[a:b]$ retrieves from index a (included) to b (excluded)
 - ♦ if a is empty $X[:b]$, the slice starts from 0
 - ♦ if b is empty $X[a:]$, the slice stops at the end of the array
 - ♦ In the case of N-dimensional arrays, each dimension can be sliced independently

Indexing and slicing (II)

- It is possible to mix integer indexing and slicing, but this will affect the rank
- Example: let `a` be an array of size (3,4). The following instructions will both copy the second row of array `a` into a new array
 - ♦ `r1 = a[1, :]`
 - ♦ `r2 = a[1:2, :]`
- However:
 - ♦ `r1` has rank 1, i.e. its shape is equal to (4)
 - ♦ `r2` has rank 2, i.e. its shape is equal to (1,4)

Reshaping

- An array can be **reshaped** by rearranging its content, as long as the total number of elements does not change

Original

(3, 4)

1	1	1	1
2	2	2	2
3	3	3	3

(6, 2)

1	1
1	1
2	2
2	2
3	3
3	3

(2, 6)

1	1	1	1	2	2
2	2	3	3	3	3

(4, 3)

1	1	1
1	2	2
2	2	3
3	3	3

(1, 12)

1	1	1	1	2	2	2	2	3	3	3	3
---	---	---	---	---	---	---	---	---	---	---	---

(12, 1)

1
1
1
1
2
2
2
2
3
3
3
3

Array math

- Basic mathematical functions operate **element-wise** on arrays, and are available both as **overloaded operators** (+, -, *, /) and as functions (add, subtract, multiply, divide)
- Most mathematical functions operate element-wise on arrays
 - ♦ E.g. `np.sqrt(x)` returns a vector, of the same shape of `x`, where each element i is the square root of x_i
- ND-array multiplications can be computed using the functions **dot** or **matmul**
 - ♦ For 1-D arrays, they return the inner product
 - ♦ For 2-D arrays, they return the matrix product
 - ♦ For N-D arrays, they have different behaviors

Vectorization – Linear regression

$$Z = Xw + b$$
$$\begin{bmatrix} z_1 \\ \vdots \\ z_m \end{bmatrix} = \begin{bmatrix} w_1 \\ \vdots \\ w_n \end{bmatrix}^T \begin{bmatrix} x_{11} & \cdots & x_{m1} \\ \vdots & & \vdots \\ x_{1n} & \cdots & x_{mn} \end{bmatrix} + \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix}$$

$$z = \text{np.dot}(w, x) + b$$

Element-wise addition

Given M examples and N features:

- x is a matrix m×n
- w and b are vectors of size n
- z is a vector of size m

2D-array multiplication

$$c_{11} = a_{11} b_{11} + a_{12} b_{21} + a_{13} b_{31} + a_{14} b_{41}$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{2e} & a_{24} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ b_{21} & b_{22} & b_{23} \\ b_{31} & b_{32} & b_{33} \\ b_{41} & b_{42} & b_{43} \end{bmatrix} = \begin{bmatrix} c_{11} & c_{12} & c_{13} \\ c_{21} & c_{22} & c_{23} \end{bmatrix}$$

2x4 **4x3** **2x3**

$$c_{22} = a_{21} b_{21} + a_{22} b_{22} + a_{23} b_{32} + a_{24} b_{42}$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{2e} & a_{24} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ b_{21} & b_{22} & b_{23} \\ b_{31} & b_{32} & b_{33} \\ b_{41} & b_{42} & b_{43} \end{bmatrix} = \begin{bmatrix} c_{11} & c_{12} & c_{13} \\ c_{21} & c_{22} & c_{23} \end{bmatrix}$$

Reminder – matrix multiplication

- Matrix multiplication is not commutative
- Matrix multiplication is not defined if the inner numbers of the sizes do not match

Array operations

- Numpy provides many useful functions for performing computations over arrays, such as taking the sum, mean, min and max of its value
- Array operations can be performed on the whole array (default) or along specific axis

axis 1
→
axis 0 ↓ $\begin{bmatrix} 1 & 2 & 3 \\ 5 & 6 & 10 \end{bmatrix}$ $\begin{bmatrix} 6 \\ 21 \end{bmatrix}$ `np.sum(a, axis=1)`

$\begin{bmatrix} 6 & 8 & 13 \end{bmatrix}$ 27 `np.sum(a)`

`np.sum(a, axis=0)`

Broadcasting

- Broadcasting allows, under certain conditions, to combine arrays with incompatible shapes during mathematical operations
- The smaller array is “broadcasted” by copying its content across the larger array, thus avoiding:
 - ♦ looping in Python (computational inefficiency)
 - ♦ creating array copies (memory inefficiency)
- Broadcasting compares the shapes of the arrays to check for their compatibility
- If the compatibility conditions are not met, a `ValueError: operands could not be broadcast together` exception is thrown

Broadcasting rules

- Broadcasting compares the shapes of the arrays starting with the trailing dimensions and going forwards
- Arrays do not need to have the same rank (number of dimension)
- Two dimensions are compatible if
 - ♦ they have the same value, or
 - ♦ one of them is 1
- The output size for each axis is the largest one (i.e., the one that is not 1)

Broadcasting rules (II)

Shapes that broadcast

A (2d) : 5 x 4
B (1d) : 1
Result (2d) : 5 x 4

A (2d) : 5 x 4
B (1d) : 4
Result (2d) : 5 x 4

A (3d) : 15 x 3 x 5
B (3d) : 15 x 1 x 5
Result (3d) : 15 x 3 x 5

A (3d) : 15 x 3 x 5
B (2d) : 3 x 5
Result (3d) : 15 x 3 x 5

A (2d) : 15 x 3 x 5
B (1d) : 3 x 1
Result (2d) : 15 x 3 x 5

Shapes that do not broadcast

A (2d) : 4
B (1d) : 3

A (2d) : 5 x 4
B (1d) : 5 x 3

A (3d) : 15 x 3 x 5
B (3d) : 15 x 2 x 5

A (3d) : 15 x 3 x 5
B (2d) : 3 x 6

A (3d) : 15 x 3 x 5
B (2d) : 4 x 1

Broadcasting example

$$\begin{bmatrix} 1 & 2 & 3 \\ 5 & 6 & 10 \end{bmatrix} * [5] = \begin{bmatrix} 1 & 2 & 3 \\ 5 & 6 & 10 \end{bmatrix} * \begin{bmatrix} 5 & 5 & 5 \\ 5 & 5 & 5 \end{bmatrix} = \begin{bmatrix} 5 & 10 & 15 \\ 25 & 30 & 50 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 5 & 6 & 10 \end{bmatrix} * [1 \ 0 \ 0] = \begin{bmatrix} 1 & 2 & 3 \\ 5 & 6 & 10 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 5 & 0 & 0 \end{bmatrix}$$

Exercise 1

- Given the amount of Carbs, Proteins, Fats in 100g of different foods:

	Pear	Broccoli	Pasta	Avocado
Carb	27	5.8	41.4	18.0
Protein	0.7	2.5	8	4.0
Fat	0.2	0.3	1.2	29.5

- knowing that carbs and proteins provides 4 calories and fats 9,
- calculate the % of calories from carbs, proteins and fats for each food.

Exercise 2

- Given two 1-D arrays of values \mathbf{x} and an array of binary labels \mathbf{y} , calculate the vector \mathbf{z} so that:

$$z_i = \begin{cases} \|x_i\|^2 & \text{if } y_i = 1 \\ \|m - x_i\|^2 & \text{if } y_i = 0 \end{cases}$$


- where m is a constant parameter.

Exercise 3

- Min-max scaling or min-max normalization is the simplest feature rescaling method and consists in rescaling each feature to the range $[0, 1]$
- The general formula for min-max normalization is as follows: $x' = \frac{x - \min(x)}{\max(x) - \min(x)}$
- Given a matrix \mathbf{x} of size M by N , where M is the number of samples and N is the number of features, write a vectorized expression to perform min-max scaling

Exercise 4

- Given 1D array, calculate the average of each consecutive triplet



References and tutorials

- <https://cs231n.github.io/python-numpy-tutorial/>
- <https://numpy.org/devdocs/user/index.html>
- <https://numpy.org/doc/stable/user/numpy-for-matlab-users.html>
- <https://numpy.org/doc/stable/reference/routines.math.html>
- <https://www.w3resource.com/python-exercises/numpy/index.php>